



第三届 eBPF开发者大会

[www.ebpftravel.com](http://www.ebpftravel.com)

# 深挖eBPF的插桩 和回调过程

张银奎 《软件调试》作者

2025-4-19

中国·西安



## 第三届 eBPF 开发者大会

[www.ebpftravel.com](http://www.ebpftravel.com)



中国·西安

```
SYSCALL_DEFINE3(bpf, int, cmd, union bpf_attr  
__user *, uattr, unsigned int, size)
```

EXPLORER

OPEN EDITORS

core.c kernel\bpf

syscall.c kernel\bpf

x

bpf.h debian/linux-headers/usr/src/linux-headers-5.10.110-r...

trace\_kprobe.c kernel\trace

kgdb.h arch\arm64\include\asm

kprobes.c arch\arm64\kernel\probes

insn.c arch\arm64\kernel

cpuidle.c drivers\cpuidle

idle.c kernel\sched

ULAN-5.10-RK3588

debian

linux-headers

usr

src\linux-headers-5.10.110-rockchip-rk3588-taiyi

include

uapi

linux

baycom.h

bcache.h

bcm933xx\_hcs.h

bfs\_fs.h

binfmts.h

blkpg.h

blktrace\_api.h

blkzoned.h

bpf\_518.h

bpf\_common\_518.h

bpf\_common.h

bpf\_perf\_event\_518.h

bpf\_perf\_event.h

bpf.h

bpfilter.h

bpqether.h

bsg.h

OUTLINE

TIMELINE

syscall.c

bpf.h

trace\_kprobe.c

kgdb.h

kprobes.c

insn.c

cpuidle.c

idle.c

debian > linux-headers > usr > src > linux-headers-5.10.110-rockchip-rk3588-taiyi > include > uapi > linux > bpf.h > bpf\_attr > \_unnamed\_struct\_6bb1\_6

479

};

480

481

#define BPF\_OBJ\_NAME\_LEN 16U

482

483

union bpf\_attr {

484

struct { /\* anonymous struct used by BPF\_MAP\_CREATE command \*/

485

\_\_u32 map\_type; /\* one of enum bpf\_map\_type \*/

486

\_\_u32 key\_size; /\* size of key in bytes \*/

487

\_\_u32 value\_size; /\* size of value in bytes \*/

488

\_\_u32 max\_entries; /\* max number of entries in a map \*/

489

\_\_u32 map\_flags; /\* BPF\_MAP\_CREATE related

490

\* flags defined above.

491

\*/

492

\_\_u32 inner\_map\_fd; /\* fd pointing to the inner map \*/

493

\_\_u32 numa\_node; /\* numa node (effective only if

494

\* BPF\_F\_NUMA\_NODE is set).

495

\*/

496

char map\_name[BPF\_OBJ\_NAME\_LEN];

497

\_\_u32 map\_ifindex; /\* ifindex of netdev to create on \*/

498

\_\_u32 btf\_fd; /\* fd pointing to a BTF type data \*/

499

\_\_u32 btf\_key\_type\_id; /\* BTF type\_id of the key \*/

500

\_\_u32 btf\_value\_type\_id; /\* BTF type\_id of the value \*/

501

\_\_u32 btf\_vmlinux\_value\_type\_id; /\* BTF type\_id of a kernel-

502

\* struct stored as the

503

\* map value

504

\*/

505

};

506

507

struct { /\* anonymous struct used by BPF\_MAP\_\*\_ELEM commands \*/

508

\_\_u32 map\_fd;

509

\_\_aligned\_u64 key;

510

union {

511

\_\_aligned\_u64 value;

512

\_\_aligned\_u64 next\_key;

TERMINAL

PORTS

PROBLEMS

OUTPUT

DEBUG CONSOLE

C/C++ Configuration W

[10/15/2024, 8:59:31 AM] Unable to resolve configuration with compilerPath: "D:\msys64\mingw64\bin\gcc.exe"

[10/15/2024, 8:59:31 AM] Unable to resolve configuration with compilerPath: "D:\msys64\mingw64\bin\gcc.exe"

Ln 492, Col 67

Tab Size: 4

UTF-8

LF

{ }

C

CODEGEEX

Win32

```
geduer@ulan:~$ cat HelloYourLand.py
#!/usr/bin/python
```

```
from bcc import BPF
```

```
prog = """
int hello(void *ctx){
    bpf_trace_printk("Here is YourLand!\n");
    return 0;
}
"""
```

```
b = BPF(text=prog)
b.attach_kprobe(event=b.get_syscall_fnname("clone"), fn_name="hello")
```

```
# header
print("%-18s %-16s %-6s %s" % ("TIME(s)", "COMM", "PID", "MESSAGE"))
while 1:
    try:
        (task, pid, cpu, flags, ts, msg) = b.trace_fields()
    except ValueError:
        continue
    print("%-18.9f %-16s %-d %s" % (ts, task, pid, msg))g
```



geduer@ulan:~\$ sudo ./HelloYourLand.py

[sudo] password for geduer:

TIME(s)	COMM	PID	MESSAGE
359.216558000	b'<...>'	2759	b'Here is YourLand!'
365.169917000	b'<...>'	2588	b'Here is YourLand!'
365.182842000	b'<...>'	2812	b'Here is YourLand!'
366.049768000	b'<...>'	1	b'Here is YourLand!'
366.328557000	b'sshd'	2588	b'Here is YourLand!'
366.349306000	b'sshd'	2816	b'Here is YourLand!'
366.355594000	b'<...>'	2817	b'Here is YourLand!'
366.367741000	b'<...>'	2819	b'Here is YourLand!'
366.375037000	b'<...>'	2820	b'Here is YourLand!'
366.391117000	b'00-header'	2820	b'Here is YourLand!'
366.399792000	b'00-header'	2820	b'Here is YourLand!'
366.407825000	b'00-header'	2820	b'Here is YourLand!'
366.418288000	b'run-parts'	2819	b'Here is YourLand!'
366.427935000	b'run-parts'	2819	b'Here is YourLand!'
366.433061000	b'run-parts'	2819	b'Here is YourLand!'
366.437419000	b'85-fwupd'	2827	b'Here is YourLand!'
366.445348000	b'run-parts'	2819	b'Here is YourLand!'
366.452698000	b'90-updates-avai'	2829	b'Here is YourLand!'
366.459509000	b'90-updates-avai'	2829	b'Here is YourLand!'
366.461085000	b'90-updates-avai'	2829	b'Here is YourLand!'
366.473640000	b'run-parts'	2819	b'Here is YourLand!'

```
bpf(BPF_PROG_LOAD, {prog_type=BPF_PROG_TYPE_SOCKET_FILTER, insn_cnt=2, insns=0x7fffca1100,
license="GPL", log_level=0, log_size=0, log_buf=NULL, kern_version=KERNEL_VERSION(0, 0, 0),
prog_flags=0, prog_name="libbpf_nametest", prog_ifindex=0,
expected_attach_type=BPF_CGROUP_INET_INGRESS, prog_btf_fd=0, func_info_rec_size=0,
func_info=NULL, func_info_cnt=0, line_info_rec_size=0, line_info=NULL, line_info_cnt=0, attach_btf_id=0,
attach_prog_fd=0, fd_array=NULL}, 148) = 4
close(4) = 0
bpf(BPF_PROG_LOAD, {prog_type=BPF_PROG_TYPE_KPROBE, insn_cnt=16, insns=0x7f91e18000,
license="GPL", log_level=0, log_size=0, log_buf=NULL, kern_version=KERNEL_VERSION(5, 10, 110),
prog_flags=0, prog_name="hello", prog_ifindex=0, expected_attach_type=BPF_CGROUP_INET_INGRESS,
prog_btf_fd=3, func_info_rec_size=8, func_info=0x48bc550, func_info_cnt=1, line_info_rec_size=16,
line_info=0x591fe30, line_info_cnt=5, attach_btf_id=0, attach_prog_fd=0, fd_array=NULL}, 148) = 4
openat(AT_FDCWD, "/sys/bus/event_source/devices/kprobe/type", O_RDONLY) = 5
read(5, "6\n", 4096) = 2
close(5) = 0
openat(AT_FDCWD, "/sys/bus/event_source/devices/kprobe/format/retprobe", O_RDONLY) = 5
read(5, "config:0\n", 4096) = 9
close(5) = 0
perf_event_open({type=0x6 /* PERF_TYPE_??? */, size=0x88 /* PERF_ATTR_SIZE_??? */, config=0,
sample_period=1, sample_type=0, read_format=0, precise_ip=0 /* arbitrary skid */, ...}, -1, 0, -1,
PERF_FLAG_FD_CLOEXEC) = 5
ioctl(5, PERF_EVENT_IOC_SET_BPF, 4) = 0
ioctl(5, PERF_EVENT_IOC_ENABLE, 0) = 0
```

	COMM		COMM	PID	MESSAGE
write(1, "TIME(s)		"..., 51	TIME(s)		
) = 51					
openat(AT_FDCWD, "/sys/kernel/debug/tracing/trace_pipe", O_RDONLY O_CLOEXEC) = 6					
fstat(6, {st_mode=S_IFREG 0444, st_size=0, ...}) = 0					
ioctl(6, TCGETS, 0x7ffca1c20)		= -1	ENOTTY (Inappropriate ioctl for device)		
lseek(6, 0, SEEK_CUR)		= -1	ESPIPE (Illegal seek)		
read(6, 0x4f75f10, 4096)		= -1	EINTR (Interrupted system call)		





## 第三届 eBPF 开发者大会

[www.ebpftravel.com](http://www.ebpftravel.com)



中国·西安

# 修补之前的\_\_arm64\_sys\_clone

```
u fffffffc0`10096100
lk! __arm64_sys_clone [kernel/fork.c @ 2643]:
ffffffc010096100 d503245f hint #0x22
ffffffc010096104 aa1e03e9 mov x9, x30
ffffffc010096108 d503201f nop
ffffffc01009610c aa0003e4 mov x4, x0
ffffffc010096110 d503233f paciasp
ffffffc010096114 a9bf7bfd stp x29, x30, [sp, #-0x10]!
ffffffc010096118 910003fd mov x29, sp
ffffffc01009611c a9400400 ldp x0, x1, [x0]
```

# 修补后的\_\_arm64\_sys\_clone

```
u fffffffc0`10096100  
lk!__arm64_sys_clone [kernel/fork.c @ 2643]:
```

```
fffffffc010096100 d4200080 brk #4  
fffffffc010096104 aa1e03e9 mov x9, x30  
fffffffc010096108 d503201f nop  
fffffffc01009610c aa0003e4 mov x4, x0  
fffffffc010096110 d503233f paciasp  
fffffffc010096114 a9bf7bfd stp x29, x30, [sp, #-0x10]!  
fffffffc010096118 910003fd mov x29, sp  
fffffffc01009611c a9400400 ldp x0, x1, [x0]
```

### C6.2.38 BRK

Breakpoint instruction. A BRK instruction generates a Breakpoint Instruction exception. The PE records the exception in `ESR_ELx`, using the EC value `0x3c`, and captures the value of the immediate argument in `ESR_ELx.ISS`.

31	30	29	28	27	26	25	24	23	22	21	20					5	4	3	2	1	0
1	1	0	1	0	1	0	0	0	0	1		imm16						0	0	0	0

#### Encoding

BRK #<imm>

```
arch > arm64 > include > asm > C brk-imm.h
```

```
/*  
 * #imm16 values used for BRK instruction generation  
 * 0x004: for installing kprobes  
 * 0x005: for installing uprobes  
 * 0x006: for kprobe software single-step  
 * Allowed values for kgdb are 0x400 - 0x7ff  
 * 0x100: for triggering a fault on purpose (reserved)  
 * 0x400: for dynamic BRK instruction  
 * 0x401: for compile time BRK instruction  
 * 0x800: kernel-mode BUG() and WARN() traps  
 * 0x9xx: tag-based KASAN trap (allowed values 0x900 - 0x9ff)  
 */
```

```
#define KPROBES_BRK_IMM 0x004  
#define UPROBES_BRK_IMM 0x005  
#define KPROBES_BRK_SS_IMM 0x006  
#define FAULT_BRK_IMM 0x100  
#define KGDB_DYN_DBG_BRK_IMM 0x400  
#define KGDB_COMPILED_DBG_BRK_IMM 0x401  
#define BUG_BRK_IMM 0x800  
#define KASAN_BRK_IMM 0x900  
#define KASAN_BRK_MASK 0x0ff
```

```
fffffffc010096100 d4200080 brk #4
```



```
static int __kprobes __aarch64_insn_write(void *addr, __le32 insn)
{
    void *waddr = addr;
    unsigned long flags = 0;
    int ret;

    raw_spin_lock_irqsave(&patch_lock, flags);
    waddr = patch_map(addr, FIX_TEXT_POKE0);

    ret = copy_to_kernel_nofault(waddr, &insn, AARCH64_INSN_SIZE);

    patch_unmap(FIX_TEXT_POKE0);
    raw_spin_unlock_irqrestore(&patch_lock, flags);

    return ret;
}
```

kn

#	Frame Base	Return Address	Call Site
00	ffffffc0`1761bad0	ffffffc0`11563f64	lk!aarch64_insn_patch_text [arch/arm64/kernel/insn.c @ 236]
01	ffffffc0`1761bb10	ffffffc0`101a8e5c	lk!arch_prepare_kprobe+0xa4 [arch/arm64/kernel/probes/kprobes.c @ 48]
02	ffffffc0`1761bb70	ffffffc0`101fd7dc	lk!register_kprobe+0x41c [kernel/kprobes.c @ 1721]
03	ffffffc0`1761bba0	ffffffc0`101ff7c0	lk!__register_trace_kprobe+0x108 [kernel/trace/trace_kprobe.c @ 514]
04	ffffffc0`1761bbd0	ffffffc0`101f1d4c	lk!create_local_trace_kprobe+0xc0 [kernel/trace/trace_kprobe.c @ 1841]
05	ffffffc0`1761bc10	ffffffc0`10258bb0	lk!perf_kprobe_init+0x8c [kernel/trace/trace_event_perf.c @ 275]
06	ffffffc0`1761bc30	ffffffc0`10259924	lk!perf_kprobe_event_init+0x50 [kernel/events/core.c @ 9619]
07	ffffffc0`1761bc60	ffffffc0`1025d2b8	lk!perf_try_init_event+0x54 [kernel/events/core.c @ 11001]
08	ffffffc0`1761bd00	ffffffc0`10264580	lk!perf_event_alloc+0x384 [kernel/events/core.c @ 11052]
09	ffffffc0`1761be20	ffffffc0`1026518c	lk!__do_sys_perf_event_open+0x180 [kernel/events/core.c @ 11808]
0a	ffffffc0`1761be30	ffffffc0`1002b668	lk!__arm64_sys_perf_event_open+0x2c [kernel/events/core.c @ 11698]
0b	ffffffc0`1761be70	ffffffc0`1002b8ec	lk!el0_svc_common.constprop.0+0xa8 [arch/arm64/kernel/syscall.c @ 36]
0c	ffffffc0`1761be80	ffffffc0`11556170	lk!do_el0_svc+0x7c [arch/arm64/kernel/syscall.c @ 195]
0d	ffffffc0`1761bea0	ffffffc0`11556c44	lk!el0_svc+0x20 [arch/arm64/kernel/entry-common.c @ 358]
0e	ffffffc0`1761beb0	ffffffc0`100127e0	lk!el0_sync_handler+0xa4 [arch/arm64/kernel/entry-common.c @ 374]
0f	ffffffc0`1761bff0	00000000`00000000	lk!el0_sync+0x1a0 [arch/arm64/kernel/entry.S @ 789]

Settings Nano主页 Nano调试

文件 视图 输出 调试 高级 帮助

ffffffffffc0`13eb3db0 ffffffff0`101969f0 lk!multi\_cpu\_stop+0xc0 [kernel/stop\_machine.c @ 220]  
ffffffffffc0`13eb3e10 ffffffff0`100cdefc lk!cpu\_stopper\_thread+0xa0 [kernel/stop\_machine.c @ 533]  
ffffffffffc0`13eb3e50 ffffffff0`100c51e0 lk!smpboot\_thread\_fn+0x15c [kernel/smpboot.c @ 142]  
ffffffffffc0`13eb3eb0 ffffffff0`10017b1c lk!kthread+0x130 [kernel/kthread.c @ 313]  
00000000`00000010 00000008`00000008 lk!ret\_from\_fork+0x10 [arch/arm64/kernel/entry.S @ 1138]  
Switched to CPU4:  
Frame Base Return Address Call Site  
ffffffffffc0`13ebdbd0 ffffffff0`10196d34 lk!aarch64\_insn\_patch\_text\_cb+0x54 [arch/arm64/kernel/insn.c @ 228]  
ffffffffffc0`13ebdbd0 ffffffff0`101969f0 lk!multi\_cpu\_stop+0xc0 [kernel/stop\_machine.c @ 220]  
ffffffffffc0`13ebde10 ffffffff0`100cdefc lk!cpu\_stopper\_thread+0xa0 [kernel/stop\_machine.c @ 533]  
ffffffffffc0`13ebde50 ffffffff0`100c51e0 lk!smpboot\_thread\_fn+0x15c [kernel/smpboot.c @ 142]  
ffffffffffc0`13ebdeb0 ffffffff0`10017b1c lk!kthread+0x130 [kernel/kthread.c @ 313]  
00000000`00000010 00000008`00000008 lk!ret\_from\_fork+0x10 [arch/arm64/kernel/entry.S @ 1138]  
Switched to CPU5:  
Frame Base Return Address Call Site  
ffffffffffc0`13f03d40 ffffffff0`10196d34 lk!aarch64\_insn\_patch\_text\_cb+0x54 [arch/arm64/kernel/insn.c @ 228]  
ffffffffffc0`13f03db0 ffffffff0`101969f0 lk!multi\_cpu\_stop+0xc0 [kernel/stop\_machine.c @ 220]  
ffffffffffc0`13f03e10 ffffffff0`100cdefc lk!cpu\_stopper\_thread+0xa0 [kernel/stop\_machine.c @ 533]  
ffffffffffc0`13f03e50 ffffffff0`100c51e0 lk!smpboot\_thread\_fn+0x15c [kernel/smpboot.c @ 142]  
ffffffffffc0`13f03eb0 ffffffff0`10017b1c lk!kthread+0x130 [kernel/kthread.c @ 313]  
00000000`00000010 00000008`00000008 lk!ret\_from\_fork+0x10 [arch/arm64/kernel/entry.S @ 1138]  
Switched to CPU6:  
Frame Base Return Address Call Site  
ffffffffffc0`13f2bd40 ffffffff0`10196d34 lk!aarch64\_insn\_patch\_text\_cb+0x54 [arch/arm64/kernel/insn.c @ 228]  
ffffffffffc0`13f2bdb0 ffffffff0`101969f0 lk!multi\_cpu\_stop+0xc0 [kernel/stop\_machine.c @ 220]  
ffffffffffc0`13f2be10 ffffffff0`100cdefc lk!cpu\_stopper\_thread+0xa0 [kernel/stop\_machine.c @ 533]  
ffffffffffc0`13f2be50 ffffffff0`100c51e0 lk!smpboot\_thread\_fn+0x15c [kernel/smpboot.c @ 142]  
ffffffffffc0`13f2beb0 ffffffff0`10017b1c lk!kthread+0x130 [kernel/kthread.c @ 313]  
00000000`00000010 00000008`00000008 lk!ret\_from\_fork+0x10 [arch/arm64/kernel/entry.S @ 1138]  
Switched to CPU7:  
Frame Base Return Address Call Site  
ffffffffffc0`13f53d10 ffffffff0`11563608 lk!\_\_aarch64\_insn\_write [arch/arm64/kernel/insn.c @ 140]  
ffffffffffc0`13f53d40 ffffffff0`10196d34 lk!aarch64\_insn\_patch\_text\_cb+0xb4 [arch/arm64/kernel/insn.c @ 158]  
ffffffffffc0`13f53db0 ffffffff0`101969f0 lk!multi\_cpu\_stop+0xc0 [kernel/stop\_machine.c @ 220]  
ffffffffffc0`13f53e10 ffffffff0`100cdefc lk!cpu\_stopper\_thread+0xa0 [kernel/stop\_machine.c @ 533]  
ffffffffffc0`13f53e50 ffffffff0`100c51e0 lk!smpboot\_thread\_fn+0x15c [kernel/smpboot.c @ 142]  
ffffffffffc0`13f53eb0 ffffffff0`10017b1c lk!kthread+0x130 [kernel/kthread.c @ 313]  
00000000`00000010 00000008`00000008 lk!ret\_from\_fork+0x10 [arch/arm64/kernel/entry.S @ 1138]  
\*BUSY\* Running...

反汇编 nd C insn.c

```
115     return (void *)set_fixmap_offset(fixmap, page_to_phys(page) +
116         (uintaddr & ~PAGE_MASK));
117 }
118
119 static void __kprobes patch_unmap(int fixmap)
120 {
121     clear_fixmap(fixmap);
122 }
123
124 /* In ARMv8-A, A64 instructions have a fixed length of 32 bits and are always
125  * little-endian.
126  */
127 int __kprobes aarch64_insn_read(void *addr, u32 *insnp)
128 {
129     int ret;
130     __le32 val;
131
132     ret = copy_from_kernel_nofault(&val, addr, AARCH64_INSN_SIZE);
133     if (!ret)
134         *insnp = le32_to_cpu(val);
135
136     return ret;
137 }
138
139 static int __kprobes __aarch64_insn_write(void *addr, __le32 insn)
140 {
141     void *waddr = addr;
142     unsigned long flags = 0;
143     int ret;
144
145     raw_spin_lock_irqsave(&patch_lock, flags);
146     waddr = patch_map(addr, FIX_TEXT_POKE0);
147
148     ret = copy_to_kernel_nofault(waddr, &insn, AARCH64_INSN_SIZE);
149
150     patch_unmap(FIX_TEXT_POKE0);
151     raw_spin_unlock_irqrestore(&patch_lock, flags);
152
153     return ret;
154 }
155
156 int __kprobes aarch64_insn_write(void *addr, u32 insn)
157 {
158     return __aarch64_insn_write(addr, cpu_to_le32(insn));
159 }
160
161 bool __kprobes aarch64_insn_uses_literal(u32 insn)
162 {
163     /* ldr/ldrsw (literal), prfm */
164
165     return aarch64_insn_is_ldr_lit(insn) ||
166         aarch64_insn_is_ldrsw_lit(insn) ||
```

```
dt lk!aarch64_insn_patch 0xffffffffc01761bab0
    +0x000 text_addrs      : 0xffffffffc0`11563928  ->
0xd5384103`a9bd7bfd
    +0x008 new_insns       : 0x00000000`d503233f  -> ??
    +0x010 insn_cnt        : 2
    +0x014 cpu_count       : 8
```

```
dd 0xffffffffc01761bab0
ffffffc0`1761bab0  11563928 ffffffc0 d503233f 00000000
ffffffc0`1761bac0  00000002 00000008 f3bd1600 f106c3c9
ffffffc0`1761bad0  1761bb10 ffffffc0 101a8e5c ffffffc0
ffffffc0`1761bae0  4fe5ca18 ffffff81 00000000 00000000
ffffffc0`1761baf0  d503245f d42000c0 13e9b000 ffffffc0
ffffffc0`1761bb00  13e9b004 ffffffc0 f3bd1600 f106c3c9
ffffffc0`1761bb10  1761bb70 ffffffc0 101fd7dc ffffffc0
ffffffc0`1761bb20  00000000 00000000 4fe5ca00 ffffff81..
```

# CPU4 Master

kn

#	Frame Base	Return Address	Call Site
00	ffffffc0`162f1900	ffffffc0`102bbd60	lk!vmacache_find+0x34 [mm/vmacache.c @ 32]
01	ffffffc0`162f1930	ffffffc0`11564c14	lk!find_vma+0x2c [mm/mmap.c @ 2411]
02	ffffffc0`162f19a0	ffffffc0`11564f40	lk!do_page_fault+0x224 [arch/arm64/mm/fault.c @ 608]
03	ffffffc0`162f19c0	ffffffc0`1003a770	lk!do_translation_fault+0xbc [arch/arm64/mm/fault.c @ 703]
04	ffffffc0`162f1a00	ffffffc0`115557b8	lk!do_mem_abort+0x4c [arch/arm64/mm/fault.c @ 837]
05	ffffffc0`162f1a40	ffffffc0`11556110	lk!e1l_abort+0x98 [arch/arm64/kernel/entry-common.c @ 123]
06	ffffffc0`162f1a50	ffffffc0`10012408	lk!e1l_sync_handler+0xb0 [arch/arm64/kernel/entry-common.c @ 123]
07	ffffffc0`162f1b90	ffffffc0`1016d330	lk!e1l_sync+0x88 [arch/arm64/kernel/entry.S @ 771]
08	ffffffc0`162f1ab0	ffff81`0b4a5c40	lk!exit_robust_list+0x80 [./arch/arm64/include/asm/uaccess.h @ 123]



Ready tasks of CPU 1:

on_cpu	state	PID	prio	wait-time	sum-exec	summ-sleep	comm
1	00000000	16	0	0	24893463	0	migration/1
0	00000000	2638	120	0	511561172	0	sshd

Ready tasks of CPU 2:

on_cpu	state	PID	prio	wait-time	sum-exec	summ-sleep	comm
1	00000000	21	0	0	30923962	0	migration/2

Ready tasks of CPU 3:

on_cpu	state	PID	prio	wait-time	sum-exec	summ-sleep	comm
1	00000000	26	0	0	23263629	0	migration/3

Ready tasks of CPU 4:

on_cpu	state	PID	prio	wait-time	sum-exec	summ-sleep	comm
1	00000000	31	0	0	4458709	0	migration/4

Ready tasks of CPU 5:

on_cpu	state	PID	prio	wait-time	sum-exec	summ-sleep	comm
1	00000000	36	0	0	5118461	0	migration/5

Ready tasks of CPU 6:

on_cpu	state	PID	prio	wait-time	sum-exec	summ-sleep	comm
1	00000000	41	0	0	3778248	0	migration/6

Ready tasks of CPU 7:

on_cpu	state	PID	prio	wait-time	sum-exec	summ-sleep	comm
1	00000000	46	0	0	2650080	0	migration/7
0	00000000	2585	120	0	3353868545	0	HelloYourLand.p

bl

0	e	ffffffc0`111b400c	e	4	0001	(0001)	lk!copy_bpf_fprog_from_user+c
1	e	ffffffc0`111b7f8c	e	4	0001	(0001)	lk!bpf_prog_create+c
2	e	ffffffc0`10201ca0	e	4	0001	(0001)	lk!bpf_jit_compile+c
3	e	ffffffc0`101ffabc	e	4	0001	(0001)	lk!bpf_patch_insn_single+c



## 第三届 eBPF 开发者大会

[www.ebpftravel.com](http://www.ebpftravel.com)



中国·西安

```
int copy_bpf_fprog_from_user(struct sock_fprog *dst, sockptr_t src, int len);
```

```
int bpf_prog_create(struct bpf_prog **pfp, struct sock_fprog_kern *fprog);
```

```
int bpf_prog_create_from_user(struct bpf_prog **pfp, struct sock_fprog *fprog,  
                             bpf_aux_classic_check_t trans, bool save_orig);
```

```
struct bpf_prog *bpf_int_jit_compile(struct bpf_prog *prog);  
void bpf_jit_compile(struct bpf_prog *prog);  
bool bpf_jit_needs_zext(void);
```

```
const char *__bpf_address_lookup(unsigned long addr, unsigned long *size,  
                                  unsigned long *off, char *sym);
```

```
bool is_bpf_text_address(unsigned long addr);
```

```
int bpf_get_kallsym(unsigned int symnum, unsigned long *value, char *type,  
                    char *sym);
```

k

Frame Base	Return Address	Call Site
ffffffc0`16bcbcb0	ffffffc0`10208268	lk!bpf_prog_load+0x8 [kernel/bpf/syscall.c @ 2123]
ffffffc0`16bcbe20	ffffffc0`10209e20	lk!__do_sys_bpf+0x6b8 [kernel/bpf/syscall.c @ 4428]
ffffffc0`16bcbe30	ffffffc0`1002b668	lk!__arm64_sys_bpf+0x2c [kernel/bpf/syscall.c @ 4383]
ffffffc0`16bcbe70	ffffffc0`1002b8ec	lk!el0_svc_common.constprop.0+0xa8 [arch/arm64/kernel/syscall.c @ 36]
ffffffc0`16bcbe80	ffffffc0`11547810	lk!do_el0_svc+0x7c [arch/arm64/kernel/syscall.c @ 195]
ffffffc0`16bcbe90	ffffffc0`115482e4	lk!enter_from_user_mode+0x10 [arch/arm64/kernel/entry-common.c @ 238]
ffffffc0`16bcbea0	0000007f`99324e44	lk!el0_sync_handler+0x74 [arch/arm64/kernel/entry-common.c @ 416]
Loading symbols for 0000007f`99300000 libbpf.so.1.2.2 -> Failed to load symbols for libbpf.so.1.2.2		
Loading symbols for 0000007f`99300000 libbpf.so.1.2.2 -> Failed to load symbols for libbpf.so.1.2.2		



```

struct bpf_prog {
    u16          pages;          /* Number of allocated pages */
    u16          jited:1,       /* Is our filter JIT'ed? */
    jit_requested:1, /* archs need to JIT the prog */
    gpl_compatible:1, /* Is filter GPL compatible? */
    cb_access:1,   /* Is control block accessed? */
    dst_needed:1,  /* Do we need dst entry? */
    blinded:1, /* Was blinded */
    is_func:1, /* program is a bpf function */
    kprobe_override:1, /* Do we override a kprobe? */
    has_callchain_buf:1, /* callchain buffer allocated? */
    enforce_expected_attach_type:1, /* Enforce expected_attach_type checking at attach time */
    call_get_stack:1; /* Do we call bpf_get_stack() or bpf_get_stackid() */
    enum bpf_prog_type type; /* Type of BPF program */
    enum bpf_attach_type expected_attach_type; /* For some prog types */
    u32 len; /* Number of filter blocks */
    u32 jited_len; /* Size of jited insns in bytes */
    u8 tag[BPF_TAG_SIZE];
    struct bpf_prog_aux *aux; /* Auxiliary fields */
    struct sock_fprog_kern *orig_prog; /* Original BPF program */
    unsigned int (*bpf_func)(const void *ctx,
                             const struct bpf_insn *insn);

    /* Instructions for interpreter */
    struct sock_filter insns[0];
    struct bpf_insn insnsi[];
};

```

```

#define BPF_PROG_RUN(prog, ctx) \
    __BPF_PROG_RUN(prog, ctx, bpf_dispatcher_nop_func)

#define __BPF_PROG_RUN(prog, ctx, dfunc) ({ \
    u32 __ret; \
    cant_migrate(); \
    if (static_branch_unlikely(&bpf_stats_enabled_key)) { \
        struct bpf_prog_stats *__stats; \
        u64 __start = sched_clock(); \
        __ret = dfunc(ctx, (prog)->insnsi, (prog)->bpf_func); \
        __stats = this_cpu_ptr(prog->aux->stats); \
        u64_stats_update_begin(&__stats->syncp); \
        __stats->cnt++; \
        __stats->nsecs += sched_clock() - __start; \
        u64_stats_update_end(&__stats->syncp); \
    } else { \
        __ret = dfunc(ctx, (prog)->insnsi, (prog)->bpf_func); \
    } \
    __ret; })

#define BPF_PROG_RUN(prog, ctx) \
    __BPF_PROG_RUN(prog, ctx, bpf_dispatcher_nop_func)

```

kn

#	Frame	Base	Return Address	Call Site
00	ffffffc0`1803bb20	ffffffc0`101fcc0c	lk!trace_call_bpf	[kernel/trace/bpf_trace.c @ 93]
01	ffffffc0`1803bbb0	ffffffc0`101fe404	lk!kprobe_perf_func+0x5c	[kernel/trace/trace_kprobe.c @ 1592]
02	ffffffc0`1803bbe0	ffffffc0`11563d68	lk!kprobe_dispatcher+0x80	[kernel/trace/trace_kprobe.c @ 1738]
03	ffffffc0`1803bc30	ffffffc0`10016cb0	lk!kprobe_breakpoint_handler+0xc8	[arch/arm64/kernel/probes/kprobes.c @ 353]
04	ffffffc0`1803bc40	ffffffc0`10017048	lk!call_break_hook+0x70	[arch/arm64/kernel/debug-monitors.c @ 326]
05	ffffffc0`1803bc60	ffffffc0`1003aa14	lk!brk_handler+0x28	[arch/arm64/kernel/debug-monitors.c @ 332]
06	ffffffc0`1803bca0	ffffffc0`115558c4	lk!do_debug_exception+0xc4	[arch/arm64/mm/fault.c @ 968]
07	ffffffc0`1803bcd0	ffffffc0`11556100	lk!el1_dbg+0x34	[arch/arm64/kernel/entry-common.c @ 186]
08	ffffffc0`1803bce0	ffffffc0`10012408	lk!el1_sync_handler+0xa0	[arch/arm64/kernel/entry-common.c @ 223]
09	ffffffc0`1803be20	ffffffc0`10095900	lk!el1_sync+0x88	[arch/arm64/kernel/entry.S @ 771]
0a	ffffffc0`1803be40	ffffffc0`1002b8ec	lk!__arm64_sys_clone	[kernel/fork.c @ 2643]
0b	ffffffc0`1803bcf0	ffffffc0`10095900	lk!do_el0_svc+0x7c	[arch/arm64/kernel/syscall.c @ 195]
0c	ffffffc0`1803bec0	0000007f`f6408b70	lk!__arm64_sys_clone	[kernel/fork.c @ 2643]





## 第三届 eBPF 开发者大会

[www.ebpftravel.com](http://www.ebpftravel.com)



中国·西安



```
int hello(void *ctx){
    bpf_trace_printk("Here is YourLand!\\n");
    return 0;
}
```

kn

#	Frame Base	Return Address	Call Site
00	ffffffc0`1803bb20	ffffffc0`101fcc0c	lk!trace_call_bpf [kernel/trace/bpf_trace.c @ 93]
01	ffffffc0`1803bbb0	ffffffc0`101fe404	lk!kprobe_perf_func+0x5c [kernel/trace/trace_kprobe.c @ 1592]
02	ffffffc0`1803bbe0	ffffffc0`11563d68	lk!kprobe_dispatcher+0x80 [kernel/trace/trace_kprobe.c @ 1738]
03	ffffffc0`1803bc30	ffffffc0`10016cb0	lk!kprobe_breakpoint_handler+0xc8 [arch/arm64/kernel/probes/kprobes.c @ 353]
04	ffffffc0`1803bc40	ffffffc0`10017048	lk!call_break_hook+0x70 [arch/arm64/kernel/debug-monitors.c @ 326]
05	ffffffc0`1803bc60	ffffffc0`1003aa14	lk!brk_handler+0x28 [arch/arm64/kernel/debug-monitors.c @ 332]
06	ffffffc0`1803bca0	ffffffc0`115558c4	lk!do_debug_exception+0xc4 [arch/arm64/mm/fault.c @ 968]
07	ffffffc0`1803bcd0	ffffffc0`11556100	lk!el1_dbg+0x34 [arch/arm64/kernel/entry-common.c @ 186]
08	ffffffc0`1803bce0	ffffffc0`10012408	lk!el1_sync_handler+0xa0 [arch/arm64/kernel/entry-common.c @ 223]
09	ffffffc0`1803be20	ffffffc0`10095900	lk!el1_sync+0x88 [arch/arm64/kernel/entry.S @ 771]
0a	ffffffc0`1803be40	ffffffc0`1002b8ec	lk!__arm64_sys_clone [kernel/fork.c @ 2643]
0b	ffffffc0`1803bcf0	ffffffc0`10095900	lk!do_el0_svc+0x7c [arch/arm64/kernel/syscall.c @ 195]
0c	ffffffc0`1803bec0	0000007f`f6408b70	lk!__arm64_sys_clone [kernel/fork.c @ 2643]

Settings Nano主页 Nano调试

文件 视图 输出 调试 高级 帮助

bp[0] was removed by ntp, total 0 instances

lk!trace\_call\_bpf:

ffffffc0101facf0 d503233f paciasep

0xffffffffc0`101facf4 struct trace\_event\_call \*trace\_event\_call\*0x00000000`d503233f\*Memory access error>unsigned int

k

Frame Base Return Address Call Site

ffffffc0`17733b20 fffffffc0`101fcc0c lk!trace\_call\_bpf [kernel/trace/bpf\_trace.c @ 93]

ffffffc0`17733bb0 fffffffc0`101fe404 lk!kprobe\_perf\_func+0x5c [kernel/trace/trace\_kprobe.c @ 1592]

ffffffc0`17733be0 fffffffc0`11563d68 lk!kprobe\_dispatcher+0x80 [kernel/trace/trace\_kprobe.c @ 1738]

ffffffc0`17733c30 fffffffc0`10016cb0 lk!kprobe\_breakpoint\_handler+0xc8 [arch/arm64/kernel/probes/kprobes.c @ 353]

ffffffc0`17733c40 fffffffc0`10017048 lk!call\_break\_hook+0x70 [arch/arm64/kernel/debug-monitors.c @ 326]

ffffffc0`17733c60 fffffffc0`1003aa14 lk!brk\_handler+0x28 [arch/arm64/kernel/debug-monitors.c @ 332]

ffffffc0`17733ca0 fffffffc0`115558c4 lk!do\_debug\_exception+0xc4 [arch/arm64/mm/fault.c @ 968]

ffffffc0`17733cd0 fffffffc0`11556100 lk!el1\_dbg+0x34 [arch/arm64/kernel/entry-common.c @ 186]

ffffffc0`17733ce0 fffffffc0`10012408 lk!el1\_sync\_handler+0xa0 [arch/arm64/kernel/entry-common.c @ 223]

ffffffc0`17733e20 fffffffc0`10095900 lk!el1\_sync+0x88 [arch/arm64/kernel/entry.S @ 771]

ffffffc0`17733e40 fffffffc0`1002b8ec lk!\_\_arm64\_sys\_clone [kernel/fork.c @ 2643]

ffffffc0`17733cf0 fffffffc0`10095900 lk!do\_el0\_svc+0x7c [arch/arm64/kernel/syscall.c @ 195]

ffffffc0`17733ec0 00000000`00000000 lk!\_\_arm64\_sys\_clone [kernel/fork.c @ 2643]

0xffffffffc0`101facf4 struct trace\_event\_call \*trace\_event\_call\*0x00000000`d503233f\*Memory access error>unsigned int

kn

# Frame Base Return Address Call Site

00 fffffffc0`17733b20 fffffffc0`101fcc0c lk!trace\_call\_bpf [kernel/trace/bpf\_trace.c @ 93]

01 fffffffc0`17733bb0 fffffffc0`101fe404 lk!kprobe\_perf\_func+0x5c [kernel/trace/trace\_kprobe.c @ 1592]

02 fffffffc0`17733be0 fffffffc0`11563d68 lk!kprobe\_dispatcher+0x80 [kernel/trace/trace\_kprobe.c @ 1738]

03 fffffffc0`17733c30 fffffffc0`10016cb0 lk!kprobe\_breakpoint\_handler+0xc8 [arch/arm64/kernel/probes/kprobes.c @ 353]

04 fffffffc0`17733c40 fffffffc0`10017048 lk!call\_break\_hook+0x70 [arch/arm64/kernel/debug-monitors.c @ 326]

05 fffffffc0`17733c60 fffffffc0`1003aa14 lk!brk\_handler+0x28 [arch/arm64/kernel/debug-monitors.c @ 332]

06 fffffffc0`17733ca0 fffffffc0`115558c4 lk!do\_debug\_exception+0xc4 [arch/arm64/mm/fault.c @ 968]

07 fffffffc0`17733cd0 fffffffc0`11556100 lk!el1\_dbg+0x34 [arch/arm64/kernel/entry-common.c @ 186]

08 fffffffc0`17733ce0 fffffffc0`10012408 lk!el1\_sync\_handler+0xa0 [arch/arm64/kernel/entry-common.c @ 223]

09 fffffffc0`17733e20 fffffffc0`10095900 lk!el1\_sync+0x88 [arch/arm64/kernel/entry.S @ 771]

0a fffffffc0`17733e40 fffffffc0`1002b8ec lk!\_\_arm64\_sys\_clone [kernel/fork.c @ 2643]

0b fffffffc0`17733cf0 fffffffc0`10095900 lk!do\_el0\_svc+0x7c [arch/arm64/kernel/syscall.c @ 195]

0c fffffffc0`17733ec0 00000000`00000000 lk!\_\_arm64\_sys\_clone [kernel/fork.c @ 2643]

0xffffffffc0`101facf4 struct trace\_event\_call \*trace\_event\_call\*0x00000000`d503233f\*Memory access error>unsigned int

6: kd>

entry.S process.c bpf\_trace.c 返回编译

86 \* Return: BPF programs always return an integer which is interpreted by

87 \* kprobe handler as:

88 \* 0 - return from kprobe (event is filtered out)

89 \* 1 - store kprobe event into ring buffer

90 \* Other values are reserved and currently alias to 1

91 \*/

92 unsigned int trace\_call\_bpf(struct trace\_event\_call \*call, void \*ctx)

93 {

94 unsigned int ret;

95

96 cant\_sleep();

97

98 if (unlikely(\_\_this\_cpu\_inc\_return(bpf\_prog\_active) != 1)) {

99 /\*

100 \* since some bpf program is already running on this cpu,

101 \* don't call into another bpf program (same or different)

102 \* and don't send kprobe event into ring-buffer,

103 \* so return zero here

104 \*/

105 ret = 0;

106 goto out;

107 }

108

109 /\*

110 \* Instead of moving rcu\_read\_lock/rcu\_dereference/rcu\_read\_unlock

111 \* to all call sites, we did a bpf\_prog\_array\_valid() there to check

112 \* whether call->prog\_array is empty or not, which is

113 \* a heuristic to speed up execution.

114 \*

115 \* If bpf\_prog\_array\_valid() fetched prog\_array was

116 \* non-NULL, we go into trace\_call\_bpf() and do the actual

117 \* proper rcu\_dereference() under RCU lock.

118 \* If it turns out that prog\_array is NULL then, we bail out.

119 \* For the opposite, if the bpf\_prog\_array\_valid() fetched pointer

120 \* was NULL, you'll skip the prog\_array with the risk of missing

121 \* out of events when it was updated in between this and the

122 \* rcu\_dereference() which is accepted risk.

123 \*/

124 ret = BPF\_PROG\_RUN\_ARRAY\_CHECK(call->prog\_array, ctx, BPF\_PROG\_RUN);

125

126 out:

127 \_\_this\_cpu\_dec(bpf\_prog\_active);

128

129 return ret;

130 }

131

132 #ifdef CONFIG\_BPF\_KPROBE\_OVERRIDE

133 BPF\_CALL\_2(bpf\_override\_return, struct trace\_event\_call \*, call, unsigned int, ret)

134 {

135 regs\_set\_return\_value(regs, rc);

136 override\_function\_with\_return(ret);

137 return 0;

Extensions have been modified on disk. Please reload the window.

Reload Window

dt call

Local var @ x0 Type trace\_event\_call\*

+0x000	list	:	list_head
+0x010	class	:	0xffffffff81`a1882408
+0x018		:	
+0x020	event	:	trace_event
+0x050	print_fmt	:	0xffffffff81`946f9280 ""(%1x)", REC->__probe_ip"
+0x058	filter	:	(null)
+0x060	mod	:	(null)
+0x068	data	:	(null)
+0x070	flags	:	32
+0x074	perf_refcount	:	1
+0x078	perf_events	:	0x00000039`d538a210
+0x080	prog_array	:	0xffffffff81`946f9c00
+0x088	perf_perm	:	(null)

r

x0=0xffffffffc01721ba58 x1=0x00000000000000013 x2=0xffffffffc00013110c x3=0x00000000000000002  
x4=0x00000000000000000 x5=0x00000000000000000 x6=0xffffffffc0101fad24 x7=0x00000000000000000  
x8=0x00000000000000007f x9=0x000000000000000001 x10=0xffffffffc0101f8d84 x11=0x00000000000000000  
x12=0x00000000000000000 x13=0x00000000000000000 x14=0x00000000000000000 x15=0x00000000000000000  
x16=0x00000000000000000 x17=0x00000000000000000 x18=0x00000000000000000 x19=0xffffffffc01477f000  
x20=0xffffffffc01721bce0 x21=0x000000000000000001 x22=0xffffffffc012c4b450 x23=0xffffffffc01721bce0  
x24=0xffffffffc013b23490 x25=0xffffffffc01721ba70 x26=0x00000000000000000 x27=0x00000000000000000  
x28=0xffffffff8138d69600 fp=0xffffffffc01721baa0 lr=0xffffffffc000131198 sp=0xffffffffc01721ba50  
pc=0xffffffffc0101f8d84 pstate=0x00000000604003c9 - Z C - - - P - D A I F EL1

db 0xffffffffc01721ba58

ffffffc0`1721ba58	48 65 72 65 20 69 73 20-59 6f 75 72 4c 61 6e 64	Here is YourLand
ffffffc0`1721ba68	21 0a 00 12 c0 ff ff ff-c8 b7 39 13 c0 ff ff ff	!.....9.....
ffffffc0`1721ba78	10 74 aa 20 81 ff ff ff-01 00 00 00 00 00 00 00	.t. ....
ffffffc0`1721ba88	50 b4 c4 12 c0 ff ff ff-00 f0 77 14 c0 ff ff ff	P.....w.....
ffffffc0`1721ba98	e0 bc 21 17 c0 ff ff ff-b0 ba 21 17 c0 ff ff ff	..!.....!.....
ffffffc0`1721baa8	9c ad 1f 10 c0 ff ff ff-20 bb 21 17 c0 ff ff ff	.....!.....
ffffffc0`1721bab8	0c cc 1f 10 c0 ff ff ff-00 59 09 10 c0 ff ff ff	.....Y.....
ffffffc0`1721bac8	e0 bc 21 17 c0 ff ff ff-e0 bc 21 17 c0 ff ff ff	..!.....!.....



```
root@ulan:~# ls -l /proc/2739/fd
```

```
total 0
```

```
lrwx----- 1 root root 64 Oct 16 17:13 0 -> /dev/pts/2
```

```
lrwx----- 1 root root 64 Oct 16 17:31 1 -> /dev/pts/2
```

```
lrwx----- 1 root root 64 Oct 16 17:31 2 -> /dev/pts/2
```

```
lr-x----- 1 root root 64 Oct 16 17:11 3 -> anon_inode:btf
```

```
lrwx----- 1 root root 64 Oct 16 17:31 4 -> anon_inode:bpf-prog
```

```
lrwx----- 1 root root 64 Oct 16 17:31 5 -> 'anon_inode:[perf_event]'
```

```
lr-x----- 1 root root 64 Oct 16 17:31 6 -> /sys/kernel/debug/tracing/trace_pipe
```





## 第三届 eBPF 开发者大会

[www.ebpftravel.com](http://www.ebpftravel.com)



中国·西安

# 调试工具1: strace

- 监视系统调用
- 从头监视
  - `strace <命令行>`
- 中途附加
  - `strace -p 2739`

# 监视收到的事件数据

```
root@ulan:~# strace -e read=6 -e trace=read -p 2739
```

```
strace: Process 2739 attached
```

```
read(6, "          bash-3109      [000] d"... , 4096) = 87
```

00000	20 20 20 20 20 20 20 20	20 20 20 20 62 61 73 68	bash
00010	2d 33 31 30 39 20 20 20	20 5b 30 30 30 5d 20 64	-3109 [000] d
00020	2e 2e 2e 20 20 31 33 35	30 2e 30 34 33 34 39 33	... 1350.043493
00030	3a 20 62 70 66 5f 74 72	61 63 65 5f 70 72 69 6e	: bpf_trace_prin
00040	74 6b 3a 20 48 65 72 65	20 69 73 20 59 6f 75 72	tk: Here is Your
00050	4c 61 6e 64 21 0a 0a		Land!..

```
read(6,
```

# 调试工具1: bpftool

- 集成在内核代码主线里的bpf工具
- tools/bpf/bpftool
- tools/lib/bpf
- 也可以使用git上的项目仓库
  - <https://github.com/libbpf/bpftool>

```
geduer@ulan:/gewu/bpftool/src$ bpftool
Usage: bpftool [OPTIONS] OBJECT { COMMAND | help }
    bpftool batch file FILE
    bpftool version
```

```
OBJECT := { prog | map | link | cgroup | perf | net | feature | btf | gen | struct_ops | iter }
OPTIONS := { {-j|--json} [{-p|--pretty}] | {-d|--debug} |
             {-V|--version} }
```

```
geduer@ulan:/gewu/ulan-5.10-rk3588/tools/bpf/bpftool$ make
```

```
make[1]: Entering directory '/gewu/ulan-5.10-rk3588/tools/lib/bpf'
```

```
Auto-detecting system features:
```

```
...      libelf: [ on ]  
...      zlib: [ on ]  
...      bpf: [ on ]
```

```
GEN      bpf_helper_defs.h  
MKDIR     staticobjs/  
CC      staticobjs/libbpf.o  
CC      staticobjs/bpf.o  
CC      staticobjs/nlattr.o  
CC      staticobjs/btf.o  
CC      staticobjs/libbpf_errno.o  
CC      staticobjs/str_error.o  
CC      staticobjs/netlink.o  
CC      staticobjs/bpf_prog_info.o  
CC      staticobjs/libbpf_probes.o  
CC      staticobjs/xsk.o  
CC      staticobjs/hashmap.o  
CC      staticobjs/btf_dump.o  
CC      staticobjs/ringbuf.o  
LD      staticobjs/libbpf-in.o  
LINK     libbpf.a  
make[1]: Leaving directory '/gewu/ulan-5.10-rk3588/tools/lib/bpf'  
LINK     bpftool
```

```
geduer@ulan:/gewu/ulan-5.10-rk3588/tools/bpf/bpftool$ sudo ./bpftool prog list
3: cgroup_device name sd_devices tag a42d275341448247 gpl
   loaded_at 2024-10-15T09:56:18+0800 uid 0
   xlated 504B jited 472B memlock 4096B
4: cgroup_skb name sd_fw_egress tag 6deef7357e7b4530 gpl
   loaded_at 2024-10-15T09:56:18+0800 uid 0
   xlated 64B jited 104B memlock 4096B
5: cgroup_skb name sd_fw_ingress tag 6deef7357e7b4530 gpl
   loaded_at 2024-10-15T09:56:18+0800 uid 0
   xlated 64B jited 104B memlock 4096B
```

```
44: kprobe name hello tag 692e7f7aa8f18bf6 gpl
    loaded_at 2024-10-15T10:48:57+0800 uid 0
    xlated 128B jited 176B memlock 4096B
    btf_id 1
```



# 调试工具3：Linux内核调试

- 虚拟机环境
  - KDB/KGDB
  - QEMU
- 真机环境
  - 挥码枪硬件调试器
  - 主机端可以是Windows或者Linux





# 格物致知

<https://nanocode.cn/>